

CSCI251/CSCI851 Autumn-2020
Advanced Programming (**SGPd**)

C++ good practice:
Part D

On we go ...

- The last set finished at the end of Part I of the textbook.
- In this set we begin with Part II, which starts on page 307.
- Part II deals with the C++ library.

Page 310-311: IO classes

- For IO processing there are 3 header files you are likely to use:
 - `iostream` : General input/output streams.
 - `fstream` : File streams.
 - `sstream` : String streams.
- The types within those are related through inheritance, so for example, `ifstream` and `istringstream` are derived from `istream`.
 - The interaction with the derived types is the same as with the base type.

- IO objects, so stream objects, don't have copy or assign.
- So you cannot do something like
`ofstream out1, out2;`
`out1=out2;`
- This also means that when we want to pass streams, we must do so by reference.
- Furthermore reading or writing an IO objects changing the state so the reference cannot be `const`.

Page 313: IO Library Condition State

Table 8.2: IO Library Condition State	
<code>strm::iostate</code>	<code>strm</code> is one of the IO types listed in Table 8.1 (p. 310). <code>iostate</code> is a machine-dependent integral type that represents the condition state of a stream.
<code>strm::badbit</code>	<code>strm::iostate</code> value used to indicate that a stream is corrupted.
<code>strm::failbit</code>	<code>strm::iostate</code> value used to indicate that an IO operation failed.
<code>strm::eofbit</code>	<code>strm::iostate</code> value used to indicate that a stream hit end-of-file.
<code>strm::goodbit</code>	<code>strm::iostate</code> value used to indicate that a stream is not in an error state. This value is guaranteed to be zero.
<code>s.eof()</code>	true if <code>eofbit</code> in the stream <code>s</code> is set.
<code>s.fail()</code>	true if <code>failbit</code> or <code>badbit</code> in the stream <code>s</code> is set.
<code>s.bad()</code>	true if <code>badbit</code> in the stream <code>s</code> is set.
<code>s.good()</code>	true if the stream <code>s</code> is in a valid state.
<code>s.clear()</code>	Reset all condition values in the stream <code>s</code> to valid state. Returns void.
<code>s.clear(flags)</code>	Reset the condition of <code>s</code> to <code>flags</code> . Type of <code>flags</code> is <code>strm::iostate</code> . Returns void.
<code>s.setstate(flags)</code>	Adds specified condition(s) to <code>s</code> . Type of <code>flags</code> is <code>strm::iostate</code> . Returns void.
<code>s.rdstate()</code>	Returns current condition of <code>s</code> as a <code>strm::iostate</code> value.

- This is a copy of Table 8.2 in the textbook.

Page 315: Buffer flushing

- Warning: Buffers aren't flushed if a program crashes.
- If you are putting output statements in for debugging you need to make sure they are flushed or you may be misled as to the crash point.
- Options for flushing:
 - `endl` : Newline then flushes.
 - `flush` : Just flushes.
 - `ends` : Null then flushes.

- Hopefully you will remember that `cerr` flushes immediately, `cout` doesn't.
- We can use a stream manipulator to change `cout` to flushing immediately.

```
cout << unitbuf;
```

- ... and turn it back again ...

```
cout << nunitbuf;
```

Pages 315-316: Tying streams

- Reading from an input stream that is tied to an output stream also results in the output stream being flushed.
 - Such tying makes sense for interactive systems.
 - The objects `cout` and `cin` are tied together.
- Tying is done using `tie`, an overloaded function with two forms:
 - Using `tie` with no argument returns a pointer to the `ostream` the object is currently tied to, or the null pointer if the stream isn't tied.
 - `x.tie(&o)` ties the stream `x` to the output stream `o`, as long as `x` isn't already tied to a stream since each stream can be tied to at most one output stream.
 - Multiple streams can be tied to the same output stream.

Page 317: C++11 file names

- We have previously seen something like:
`ifstream in(ifile);`
- ... with `ifile` the name of a file.
- Pre C++11 `ifile` had to be a C-string, so a C-style character array.
- From C++11 on, `ifile` can be a `string`.

Page 318: `fstream` → `close`

- When an `fstream` object is destroyed, such as when it goes out of scope, `close` is called automatically on the file it is tied to.
- Here goes code for processing a list of files, with the file being closed each loop when input goes out of scope.

```
string line;
for (auto p = argv+1; p!= argv +argc;++p) {
    ifstream input(*p);
    cout << endl;
    if (input){
        getline(input,line);
        cout << string(*p) << " : 1st line :" << endl;
        cout << line << endl;
    }
    else
        cerr << "couldn't open " << string(*p) << endl;
}
cout << endl;
```

Page 319-320: File modes

- If you want to preserve existing data in a file, open the `ofstream` in either `app` or `in` mode explicitly.
- When `open` is used, the mode is always set.
 - If it's not explicit, it's implicit with the default used.
 - `ifstream`: default `in`.
 - `ofstream`: default `out`.
 - `fstream`: default `in and out`.

Page 326-328: Sequential container overview

- The default container is `vector`.
- There are some guidelines in the lecture notes.
- One important point is that if you are unsure...
 - Write your code using only operations common to both vectors and lists, so use iterators not subscripts and avoid random access.
 - These lets you easily go to list or vector depending on what you end up deciding.

Pages 331- 334:

Iterators and Containers

- The iterator range, as specified by a pair of iterators referencing elements of the same container, is pivotal in the use of containers and iterators.
 - It's left inclusive: `[begin, end)` ...
... so `end` points to one past the last element.
- When you don't need write access, you `cbegin` and `cend`.
 - Here the `c` is for `const`.

Pages 334-336:

Defining and initialising a container

- When initialising a container as a direct copy of another one, the container type and element types have to match.

- So the following works ...

```
list<string> authors={"Alice", "Bob"};  
list<string> list2(authors);
```

- But the following wouldn't...

```
deque<string> dq2(authors);
```

- Note the list initialization for `authors`, allowed from C++11.

to `vector<T>`: constructors

- There are a fair few options, listed without the allocator here.
- **Default, with an empty vector.**

```
vector<int> v0;
```

- **Initialisation with values ...**

```
vector<int> v1(10, -1); // 10 ints set to -1
```

- **Initialisation without values ...**

```
vector<int> v2(10); // 10 ints, default set to 0.
```

- **By copying from another suitable `vector<T>`.**

```
vector<int> v3(v2);
```

```
vector<int> v4=v2; // equiv. to above
```

- **Iterator based construction:**

```
vector<int> v5(v4.begin(), v4.end());
```

- **List initialisation, from C++11.**

```
vector<string> words = {"one",  
"two", "red", "blue"};
```

```
vector<int> numbers{1, 2, 3, 4, 5, 6, 7};
```


Pages 340-341:

Comparing container

- A container is compared by an element by element comparison.
 - Equal iff same length and elements all the same pairwise.
 - The comparison is based on comparing the first unequal elements of the container.
 - If the containers differ in size but every element of the shorter is equal to the corresponding element of the larger, the shorter is deemed less.
- You can only use a relational operator, $>$, $>=$, $<$, $<=$, if that operator is defined for the element type the container is storing.

Pages 341-349:

Sequential Container Operations

- Container element are copies:
 - When an object is used to initialize a container, or insert something into the container, it's a copy in the container and relationship is gone.
 - This means changing the element in the container doesn't effect the original object, and vice versa.
 - Using `emplace` passes values to a constructor that's used to make an object in the container directly.
- Insertions for `vector`, `string`, or `deque`.
 - Existing iterators, references and pointers into the container may be invalidated.
 - Deletions will call problems too, as would other resizing operations.
 - Legal to insert anywhere anyway, but it could be expensive.

- Calling `front` or `back` on an empty container is a bad thing!
- STL container member functions that remove elements don't check their argument(s).

Page 353-355: Warnings re: Iterators

- If an iterator (or pointer or reference) has been invalidated, such as through a resizing operator, using it will result in a run-time error.
 - To avoid this happening it's a good idea to minimize the region in which the iterator needs to be used.
 - In particular when you are adding or removing elements, the iterator returned by `end` is usually invalidated so in loops that add/remove elements recalls to `end` are needed, not a stored iterator.

Page 356-359: Reserving memory

- It's possible to reserve memory in some containers.
- This doesn't change the number of elements stored, so the `size()` will be unchanged, but does change how many elements can be stored before space needs to be reallocated.
 - `capacity()` and `reserve()` work for `string`, `vector`.

```
vector<int> vInt;  
vInt.push_back(1);  
cout << vInt.capacity() << " " << vInt.size() << endl;  
vInt.reserve(100);  
cout << vInt.capacity() << " " << vInt.size() << endl;
```

- The C++11 function `shrink_to_fit()` requests a reduction in the capacity to equal the size.
 - But the implementation but choose to ignore this.
- Implementations may allocate different sizes.
- The function `resize()` requests a change in size, not capacity, meaning elements are initialized.

Pages 360-368 some string things ...

- Lots of different ways to construct strings.
- Operation: `s.substr(pos, n)`
 - `n` characters starting from `pos`.
- Other operations: `assign`, `insert`, `erase`, `append`, `replace`.
 - Various overloaded forms of these, i.e. different argument sets.

- **Searching operations:** `find`, `rfind`,
`find_first_of`, `find_last_of`,
`find_first_not_of`, `find_last_not_of`
- **The function `compare` can be used to check if two strings are equal.**
 - Returns +ve, 0 (equal to), -ve.
- **C++11 Conversions:** `to_string`, `stoi`, `stol`,
`stoul`, `stoll`, `stoull`, `stof`, `stod`, `stold`.
 - Failure to convert to a number results in an `invalid_argument` exception being thrown.
 - If it's a value that cannot be represented, `out_of_range` is thrown.

Generic algorithms

- Page 378: STL Algorithms never execute container operations, they operate in terms of iterators and iterator operations.
 - They never directly change the size of underlying containers, including adding or removing elements.
- Page 379: Best to use `cbegin()` and `cend()` with read-only algorithms.
 - Use `begin()` and `end()` if the plan is to use a returned iterator to change an element's value, such as in `find`.

```
alg(beg, end, other args);  
alg(beg, end, dest, other args);  
alg(beg, end, beg2, other args);  
alg(beg, end, beg2, end2, other args);
```

- Page 380: Those algorithms that take a single iterator for a 2nd sequence, so the 2nd or 3rd form above, **assume** that the 2nd sequence is at least as large as the first.
- Page 381: The sequences referenced through the iterators may refer to different kinds of container.
- Page 382: Algorithms that write to a destination iterator **assume** that destination is large enough to hold the number of elements that are going to be written.

Page 386: Passing a function to an algorithm

- Under the other args, we can use a predicate in place of `<`, for example.
- **Predicates:** Expression that can be called and that return a value that can be used as a condition.
 - Unary: Take a single parameter.
 - Binary: Take two parameters.
- To sort words by size ... the `stable_sort` maintains alphabetical order within the same length words...

```
bool isShorter (const string &s1, const string &s2)
{   return s1.size() < s2.size();           }
```

```
sort (words.begin(), words.end(), isShorter);
stable_sort (words.begin(), words.end(), isShorter);
```

Pages 387-401: Lambdas

- Lambda expressions are callable functions, so can be used as predicates.

- The general syntax:

```
[capture list] (parameter list) -> return type { function body }
```

- Must include capture list and the function body, although the capture list may be empty:

```
auto f = [] { return 42; };  
cout << f() << endl;
```

- Page 389: Lambdas with function bodies containing anything other than a single return statement that do not specify a return type return `void`.

- Page 390: A lambda being used in a function can only use a variable local to its surrounding function if the lambda captures that variable.

- So this is okay ...

```
[size] (const string &a)
    { return a.size() >= size; };
```

- ... but this isn't ...

```
[] (const string &a)
    { return a.size() >= size; };
```

- ... since `size` wasn't captured.

- Page 391: The capture list is used for local non `static` variables, lambdas can use local `static` variables and variables declared outside the function directly.
 - In particular this means something like `cout` doesn't need to be passed as a variable.
- Page 393: When we capture a variable by reference, we must ensure that the variables exist at the time the lambda executes.
- Page 394: Keep your lambda captures simple!

Pages 397-401: Bind

- The function `bind` is new to C++11.
- Can be thought of “a general-purpose function adaptor”.
- Takes a callable object ...
- ... generates a new callable that adapts the parameter list of the original object.

```
auto newCallable = bind(callable, arg_list);
```

- The argument list contains placeholders for the arguments the `newCallable` needs.

Pages 401-402: Insert Iterators

- The three types of insert iterator are created using `back_inserter`, `front_inserter`, **or** `inserter`.
- We can only use `front_inserter` if the container has `push_front`, since that's what a `front_inserter` uses.
- We can only use `back_inserter` if the container has `push_back`, since that's what a `back_inserter` uses.

Pages 407-409: Reverse iterators

- Possibly add something here later.
- Not too important at the moment...

Back to generic algorithms

- Page 411: Many compilers won't complain if we pass the wrong category of iterator to an algorithm ☹️
- Page 415: Container specific algorithms: Use the list member versions in preference to the generic algorithms available for `lists` ***and*** `forward_lists`.